

Hermes Agents vs LangChain Agents

Written against what this repo actually builds — a trade-compliance Researcher with real tools — and what the same system would look like in LangChain. Hermes claims cite the installed runtime (v0.18.2, MIT, Nous Research); LangChain claims verified against current docs, mid-2026. Token and reliability figures are measured on this repo, not estimated.

Hermes is a runtime you configure. LangChain is a library you assemble.

Hermes hands you a finished agent and asks what it should be allowed to do. LangChain hands you the parts and asks what you want to build. Both get you a tool-using agent quickly; they differ in everything *around* the agent.

THE THREE DIMENSIONS

	HERMES	LANGCHAIN
Architecture	A running process. One fixed think-act-observe loop (<code>AIAgent.run_conversation()</code>) serves the CLI, the web gateway, cron jobs and chat platforms alike. You choose what it can reach, not how it thinks. Less to build and less to get wrong — but no way to express a flow the loop does not already support.	A library you import. <code>create_agent</code> is a shortcut over <code>LangGraph</code>; underneath you draw the graph yourself — nodes, branches, loops, pause-for-a-human. Expresses anything you can draw. You maintain all of it.
Tool definition	A small MCP server named in <code>config.yaml</code> , or a plugin taking a plain Python function. The description the model reads is generated from the function itself, so the two cannot drift apart. Hermes wraps it in injection scanning, a malware check and a stripped environment. A tool lives outside the framework and costs nothing to anyone who does not switch it on.	The <code>@tool</code> decorator, attached per agent with <code>.bind_tools()</code> . Same idea — the docstring teaches the model when to call it. MCP works too, via <code>langchain-mcp-adapters</code> , as an extra package rather than a config key. The quickest possible way to add a tool. Any sandboxing is yours to add.
State management	Four things, all included, no extra service: <code>MEMORY.md</code> facts pasted into the prompt each session; every past conversation searchable in SQLite (<code>session_search</code> , no model call involved); reusable Skills; and one optional hosted provider behind a single config key. Memory for free, on its terms — exact-word search, not meaning-based, unless you plug in a provider.	Two systems you choose and wire: a checkpointer for the current conversation (<code>SqliteSaver</code> / <code>PostgresSaver</code>), and langgraph.store for anything longer-lived — usually reached through a tool you write yourself. Nothing by default, everything possible — including meaning-based search. More wiring, more reach.

WHAT THAT MEANS IN PRACTICE

	HERMES	LANGCHAIN
The agent loop	included	included
Conversations that survive a restart	included	wire a checkpointer
Memory that outlives the session	included	wire a store + a tool to write with
A second agent	one config line (<code>delegate_task</code>)	build the graph
Swap the model	a four-line YAML file	a code change + a pip install
Web API with streaming	included	write a server
Where the agent lives	it is the process	sits inside your app

For a two-tool demo, LangChain is less work. This exercise asked for memory that outlives the session, a second agent and a model swap — which is exactly where the columns flip.

WHICH ONE FOR WHICH JOB

IF YOU ARE BUILDING	REACH FOR	BECAUSE
An internal assistant a team talks to all day — this compliance desk, an ops helper, a research aide	Hermes	It has to remember last week, run on your own hardware, and be reachable from Slack and a browser. All included; none of it is your code.
A scheduled agent — check something each morning, post the result to a channel	Hermes	Cron, messaging and sessions ship inside the same runtime. In LangChain each is a separate thing you host.
Anything that must stay on your own machines — regulated data, no third-party service	Hermes	MIT, fully local, no external dependency by design.
A feature inside a product you already have — a "summarise this ticket" button in your SaaS	LangChain	You need a library your app calls, not a second process running beside it.
Question-answering over your own documents	LangChain	Retrieval and vector stores are its home ground. Hermes does not model them at all.
A flow with approvals or branches — pause for a human, take a different path, loop until a check passes	LangChain	Hermes has one loop shape. If your process does not fit it, you cannot express it.

TWO THINGS WORTH SAYING OUT LOUD

Neither framework's value is the loop. Think-act-observe is about fifteen lines and you can write it yourself. The value is everything around it that you only discover you need once real users arrive — retries, trimming long conversations, saving sessions, searching them, deciding what is worth remembering, handing work to a second agent.

A framework's defaults are the real cost. Hermes switches on 49 tools by default, tuned for a general assistant. Cutting this agent down to the five it actually uses took the prompt from **21,373 to 6,694 tokens per call** — and tool-calling got *more* reliable, because a model choosing between five options picks better than one choosing between 49.

CHOOSING

Hermes when the agent *is* the product, and its shape fits one loop with delegation.

LangChain when the agent is one feature inside something bigger, or you need document retrieval or a flow that branches.

Neither when you have two tools, one user and nothing to remember. That is the loop plus your own functions, and a framework will cost you more in defaults than it saves you in code.

Both are MIT, and neither is settled ground. LangChain has changed its agent API twice in about a year (`AgentExecutor` → `create_react_agent` → `create_agent`), and Hermes is pre-1.0 — building this found its documentation contradicting its own source in three places. **Pin your versions and read the source, whichever you pick.**

Evidence — the fifteen-line loop written out, what a tool description actually contains, the three ways into Hermes and why this repo took the third, and where open-weight models break — is in [hermes-vs-langchain-detail.md](#).